Using project sitemaps

Table of contents

1. Introduction

With Forrest 0.6 it is now possible for projects to "plugin" to our sitemaps, without needing to copy the main sitemaps and keep them them synchonised. This will enable hassle-free update to future Forrest versions.

Note:

We advise you to spend time to understand the Apache Cocoon sitemap. See Cocoon sitemap and Cocoon concepts and related component documentation. The Forrest sitemap is broken into multiple files. The main one is **sitemap.xmap** which delegates to others. See the Sitemap Reference for a tour of the default sitemap.

2. How does it work?

If a project has a `sitemap.xmap` file in it's documentation dir, that gets mounted automatically by Forrest and becomes part of the processing: it is a preprocessing step, and is the first one to handle the request. Because of this it can serve any file directly. If it does not want to serve a file, it can simply not match the URL and Forrest will take care of it as usual.

The cool thing is that if that pipeline serves an xml representation, Forrest will provide a skinned version of it.

So if the project sitemap matches test.xml and transforms that to a correctly structured Forrest intermediate "document-v*", then the user will see test.html fully rendered by Forrest.

Of course, to resolve the directories in your sitemap it is important to use the 'project:' and 'forrest:' variables to prevent any possible issue in the future.

3. Example uses of this technique

3.1. Adding a new content type

See the section "Advanced customizations: sitemap.xmap" in the Using Forrest document and then follow the Example: Adding a new content type.